

Talk-to-Data Delivery Blueprint — Phase 7 Annexes: Controlled Pilot & User Testing

Daniel Brule · [linkedin.com/in/danielbrule](https://www.linkedin.com/in/danielbrule)
v1.0 · June 2026

- 1 How to use this annex
- 2 Activity 1: Confirm pilot operating plan, cohort and control boundary
 - 2.1 Pilot start confirmation checklist
 - 2.2 Pilot operating plan template
 - 2.3 Pilot boundary statement template
 - 2.4 Pilot governance and decision route
 - 2.5 Pilot start meeting agenda
 - 2.6 Quick readiness decision record
 - 2.7 Common mistakes to avoid
- 3 Activity 2: Prepare user, support and operator guidance
 - 3.1 Documentation readiness checklist
 - 3.2 Pilot user quick-start guide template
 - 3.3 Example user guidance — illustrative
 - 3.4 Supported and unsupported question examples
 - 3.5 Support guide template
 - 3.6 Operator and technical documentation checklist
 - 3.7 Usage and cost guidance template
 - 3.8 Onboarding session agenda
 - 3.9 Documentation gap log
- 4 Activity 3: Test pilot access, onboarding, role changes and offboarding
 - 4.1 Pilot access lifecycle test plan
 - 4.2 Access-change timing record
 - 4.3 Access boundary test matrix
 - 4.4 Traceability evidence checklist
 - 4.5 Example access lifecycle test — illustrative
 - 4.6 Access exception and workaround log
 - 4.7 Minimum evidence by delivery intent
 - 4.8 Common mistakes to avoid
- 5 Activity 4: Run controlled pilot usage and monitored sessions
 - 5.1 Pilot usage plan
 - 5.2 User cohort realism check
 - 5.3 Monitored session plan
 - 5.4 Real-use observation template
 - 5.5 Outside-scope question log
 - 5.6 Pilot behaviour evidence log
 - 5.7 Facilitator guidance for monitored sessions
 - 5.8 Evidence caveat statement template
 - 5.9 Session summary template
 - 5.10 Common mistakes to avoid
- 6 Activity 5: Capture feedback and monitor pilot signals
 - 6.1 Feedback and monitoring routes
 - 6.2 Feedback-to-trace and monitoring register
 - 6.3 Pilot monitoring dashboard structure
 - 6.4 Feedback quality checklist
 - 6.5 Structured feedback form template
 - 6.6 Usage-pattern and distinctive-user follow-up list
 - 6.7 Trust and responsible-use indicators
 - 6.8 Error, refusal and clarification log
 - 6.9 Cost and latency review template
 - 6.10 High-cost or high-latency interaction review
 - 6.11 Positive-value evidence template
 - 6.12 Feedback and monitoring diagnosis guide
 - 6.13 Feedback and monitoring review cadence
 - 6.14 Pilot signal summary template
 - 6.15 Common mistakes to avoid
- 7 Activity 6: Review answer quality, caveat understanding and user interpretation

- 7.1 Real-interaction answer-quality review template
- 7.2 Answer-quality scoring guide
- 7.3 Caveat and limitation review
- 7.4 Misunderstood-answer log
- 7.5 Follow-up interview trigger list
- 7.6 User interpretation interview guide
- 7.7 Refusal and clarification understanding review
- 7.8 Answer-format improvement log
- 7.9 Real-interaction review summary
- 7.10 Common mistakes to avoid
- 8 Activity 7: Triage issues, manage backlog and controlled change
 - 8.1 Pilot issue triage and backlog log
 - 8.2 Severity and urgency guide
 - 8.3 Decision route matrix
 - 8.4 Root-cause classification guide
 - 8.5 Continue / constrain / pause / restart decision record
 - 8.6 Hot-fix and change classification guide
 - 8.7 Hot-fix decision checklist
 - 8.8 Controlled change / release record
 - 8.9 Regression check template
 - 8.10 User and support change note template
 - 8.11 Pilot change log
 - 8.12 Phase 8 carry-forward and backlog separation
 - 8.13 Triage and change review agenda
 - 8.14 Example triage and change cases — illustrative
 - 8.15 Common mistakes to avoid
- 9 Activity 8: Assess adoption, workflow fit, supportability and Phase 8 implications
 - 9.1 Pilot decision summary template
 - 9.2 Adoption evidence review
 - 9.3 Workflow-fit assessment
 - 9.4 Value evidence template
 - 9.5 Supportability assessment
 - 9.6 Phase 8 production-readiness backlog
 - 9.7 Decision recommendation guide
 - 9.8 Decision wording examples — illustrative
 - 9.9 Pilot caveat statement template
 - 9.10 Common mistakes to avoid

1 How to use this annex

These annexes provide optional working templates for Phase 7: Controlled pilot and user testing. They are intended as practical starter materials for planning, running, monitoring and assessing a controlled Talk-to-Data pilot.

They are not all mandatory outputs. Teams should select the artefacts that match the delivery intent, risk level, user population, data sensitivity and maturity of the MVP.

For a small POC or expert-user pilot, several templates may be simplified or combined. A one-page user guide, lightweight operating plan and simple feedback log may be enough.

For a broader business pilot, decision-critical use case or MVP on a production path, the evidence standard should be higher. User documentation, access lifecycle evidence, feedback-to-trace linkage, monitoring, issue triage, controlled change, support handover and pilot decision records should be treated more seriously.

The annexes should not turn Phase 7 into a heavy governance exercise. Their role is to make the pilot controlled, observable and decision-useful. Teams should use the templates that help them answer the core Phase 7 question: whether the MVP can be used safely, usefully and supportably by real users, and what must be true before production readiness can be assessed.

2 Activity 1: Confirm pilot operating plan, cohort and control boundary

2.1 Pilot start confirmation checklist

Area	Confirmation question	Status	Owner	Evidence / note
Pilot objective	Is the pilot learning objective clear?			
User cohort	Are pilot users and groups approved?			

Access groups	Are roles, permissions and access routes confirmed?			
Supported questions	Are supported question types clear?			
Excluded questions	Are excluded / unsupported question types documented?			
Data scope	Are approved data assets and caveats confirmed?			
Answer types	Are allowed answer formats defined?			
Known limitations	Are limitations visible to users and support teams?			
Monitoring	Are usage, quality, cost, latency and risk signals available?			
Feedback route	Is feedback captured in a structured way and linkable to traces?			
Support route	Is first-line and escalation support clear?			
Pause conditions	Are stop / narrow / pause triggers explicit?			
Change route	Is approval for quick fixes and material changes defined?			
Decision forum	Is the pilot review and final decision route agreed?			

Suggested status values: Ready / Ready with caveat / Gap / Blocked / Not applicable.

2.2 Pilot operating plan template

Item	Description
Pilot name	Name of the controlled pilot.
Pilot objective	What the pilot is intended to learn or prove.
MVP version	Version of the validated MVP entering pilot.
Environment	Pilot / staging / controlled internal environment.
Pilot start date	Planned start date.
Planned review cadence	Daily, twice weekly, weekly or milestone-based.
Planned pilot duration	Directional duration and review points.
Pilot users	Named users or approved user groups.
Excluded users	Users or groups explicitly excluded.
Supported questions	Question types approved for pilot answering.
Unsupported questions	Question types the system should refuse, clarify or redirect.
Data assets in scope	Approved tables, views, semantic assets or APIs.
Answer types in scope	Number, table, chart, short narrative, caveated answer, clarification, refusal.
Known limitations	Caveats users and support teams must understand.
Usage limits	Query limits, concurrency limits, time windows or cost limits.
Monitoring owner	Person or team reviewing pilot signals.
Support owner	Person or team handling user support.
Escalation route	Route for incidents, access issues, misleading answers or exposure concerns.
Change approver	Person or forum approving quick fixes and material changes.
Pause authority	Person or forum empowered to pause, narrow or restart pilot usage.
Final decision owner	Person or forum deciding Phase 7 outcome.

2.3 Pilot boundary statement template

Purpose: create a short, plain-language statement of what the pilot covers and does not cover.

Boundary area	Approved pilot boundary
Users	
Business domain	
Data assets	
Supported questions	
Unsupported questions	
Supported answer types	
Caveats / limitations	
Access restrictions	
Usage limits	
Monitoring requirements	

Monitoring requirements	
Escalation conditions	
Pause conditions	

Example — illustrative:

Boundary area	Approved pilot boundary
Users	15 regional sales managers and 2 sales operations analysts.
Business domain	Sales performance and pipeline conversion.
Data assets	sales_revenue_mart, region_dimension, pipeline_summary_view.
Supported questions	Revenue by region, customer count by segment, pipeline conversion, month-on-month variance.
Unsupported questions	Margin, individual salesperson performance, forecasting, customer profitability.
Supported answer types	Number, table, short narrative, clarification, refusal.
Caveats / limitations	Current-month revenue is provisional; pipeline is refreshed daily, not real time.
Access restrictions	Users can only see their approved region.
Usage limits	Internal pilot use only; no export for board, regulatory or external reporting.
Monitoring requirements	Weekly review of usage, errors, refusals, latency, cost and support issues.
Escalation conditions	Possible data exposure, materially misleading answer, repeated unsupported use.
Pause conditions	Incorrect access enforcement, misleading answer in supported scope, audit trace failure.

2.4 Pilot governance and decision route

Decision type	Decision owner	Evidence required	Typical timing
Start pilot	Product owner / sponsor	Phase 6 approval, pilot plan, user list, monitoring and support route	Before pilot start
Add pilot user	Product owner + access owner	Approved cohort, access check, onboarding confirmation	During pilot
Release quick fix	Product / technical owner	Fix note, test evidence, communication note	During pilot
Release material change	Decision forum / risk owner	Retest evidence, impact assessment, boundary update	Review point
Narrow pilot	Product owner + risk owner	Issue evidence, affected users/questions/data	As needed
Pause affected scope	Pause authority	Severity evidence, risk rationale, restart condition	Immediate where needed
Restart affected scope	Pause authority	Fix evidence, retest evidence, communication note	After remediation
Proceed to Phase 8	Decision forum	Pilot learning pack, backlog, risk and supportability evidence	End of pilot
Stop pilot	Sponsor / decision forum	Value, risk, supportability or feasibility evidence	As needed

2.5 Pilot start meeting agenda

Purpose: confirm that the pilot can start without reopening earlier phase decisions.

- Confirm pilot objective and expected learning.
- Confirm approved users, access groups and excluded users.
- Confirm supported and unsupported questions.
- Confirm data assets, caveats and known limitations.
- Confirm user documentation and onboarding plan.
- Confirm support, monitoring and escalation route.
- Confirm quick-fix and material-change approval route.
- Confirm pause, narrow and restart authority.
- Confirm first review date and decision cadence.
- Confirm go / no-go to start controlled pilot.

2.6 Quick readiness decision record

Item	Decision
------	----------

Pilot start decision	Start / Start with caveats / Do not start
Main caveats	
Required actions before start	
Accepted constraints	
First review point	
Decision owner	
Date	

Example decision wording — illustrative:

The Phase 7 pilot may start for the approved regional sales cohort, limited to sales performance and pipeline questions. Margin, forecasting and salesperson-level questions remain out of scope. Current-month revenue caveats must be visible to users. Any access issue, misleading supported answer or missing trace should trigger immediate review and potential pause of the affected scope.

2.7 Common mistakes to avoid

Mistake	Why it matters
Reopening the whole scope	Activity 1 becomes a new planning phase instead of a readiness confirmation.
Starting with “friendly users” but no boundary	Positive feedback becomes hard to interpret and risky behaviour may be normalised.
Treating support as optional	The pilot becomes dependent on the build team and does not test operability.
Allowing informal quick fixes	Behaviour changes without traceability, retesting or user communication.
Not defining pause authority	Everyone sees the issue, but nobody is empowered to stop or narrow exposure.
Missing unsupported-question handling	Natural user demand expands scope without approval.

3 Activity 2: Prepare user, support and operator guidance

3.1 Documentation readiness checklist

Documentation area	Confirmation question	Minimum for small pilot	Higher standard for larger pilot	Status	Owner
Pilot access	Do users know how to access the pilot?	Link / access route shared	Access guide with screenshots and troubleshooting		
Connection / login	Is authentication clear?	Basic login instruction	SSO, access group, MFA and issue route documented		
UI basics	Do users know how to use the interface?	Short quick-start note	UI guide with examples and known limitations		
Supported questions	Are included question types clear?	Short list of examples	Supported-question catalogue by domain / answer type		
Unsupported questions	Are exclusions clear?	Short list of examples	Exclusion guide with refusal / redirect examples		
Caveats	Do users understand limitations?	Key caveats visible	Caveat guide with examples of interpretation risk		
Responsible use	Do users know how not to misuse the pilot?	Short do / don't list	Responsible-use guidance and acknowledgement		
Feedback	Is the feedback route clear?	Feedback button / form	Structured feedback linked to traces and issue categories		
Escalation	Do users know when to escalate?	Named contact / channel	Escalation route by issue severity		
Usage / cost limits	Are users aware of sensible usage expectations?	Simple usage rules	Usage thresholds, cost guidance and heavy-use escalation		
Support triage	Can support handle common issues?	Short support note	Support playbook with issue categories and escalation		
Operator guidance	Can operators diagnose pilot behaviour?	Known links to logs / traces	Operator runbook with monitoring, logs, pause and restart route		
Technical documentation	Is technical documentation usable for pilot support?	Existing Phase 5/6 docs referenced	Updated technical pack with known limitations and release notes		

Suggested status values: Ready / Ready with caveat / Gap / Blocked / Not applicable.

3.2 Pilot user quick-start guide template

--	--

Section	Content to provide
What the pilot is for	Short description of the pilot objective and intended use.
Who can use it	Approved users or user groups.
How to access it	Link, access route, login method, browser / app requirements, connection steps.
How to use the UI	Basic steps: ask a question, refine, review sources / caveats, give feedback.
What you can ask	Supported domains, metrics, question types and examples.
What is out of scope	Unsupported questions, excluded data, excluded answer types.
How to read answers	How to interpret numbers, tables, caveats, sources, refusals and clarifications.
How to use it responsibly	When to verify, when not to rely on the answer, what not to use it for.
Usage expectations	Expected usage, discouraged heavy use, known latency or cost constraints.
How to give feedback	Feedback button, form, channel, what good feedback should include.
How to escalate	When and how to report access issues, wrong answers, missing data or incidents.
Known limitations	Current limitations users should understand before using the pilot.

3.3 Example user guidance — illustrative

What this pilot is for: This pilot helps approved regional sales managers ask questions about sales performance and pipeline conversion using the approved sales data assets.

How to access the pilot: Use the pilot link shared by the product owner. Access is restricted to approved pilot users. If you cannot connect or see the expected region, contact the pilot support channel.

Good questions:

- “What was revenue last month by region?”
- “Show pipeline conversion by segment for Q2.”
- “How did customer count change month on month?”

Unsupported questions:

- Margin or profitability questions.
- Individual salesperson performance.
- Forecasting or future revenue predictions.

Questions outside the approved sales performance data scope.

- How to interpret answers
Review caveats before using an answer. Current-month revenue may be provisional. Pipeline data is refreshed daily, not in real time. If an answer looks surprising, ask a follow-up, check the caveat or escalate.

Responsible use: Do not use pilot outputs for external reporting, board packs, regulatory submissions or contractual decisions. The pilot is intended to test controlled internal use.

Usage and cost expectations: Use the pilot for real workflow questions, not high-volume exploration or repeated broad queries. Heavy or unusual usage may be reviewed if it creates cost, latency or support issues.

Feedback: Use the feedback option after useful, wrong, confusing or unsupported answers. Feedback should describe what you expected and whether the answer helped the workflow.

3.4 Supported and unsupported question examples

Question	Expected system behaviour	User guidance note
“What was revenue last month by region?”	Answer if within approved scope and user access.	Supported question. Review caveats.
“Show revenue for my region this quarter.”	Answer using user access boundary.	Supported if region access is correctly applied.
“What is margin by product?”	Refuse or redirect if margin is out of scope.	Unsupported for current pilot.
“Which salesperson performed worst?”	Refuse if individual performance is excluded.	Not supported due to scope and sensitivity.
“Forecast next quarter revenue.”	Refuse or clarify if forecasting is excluded.	Not validated for pilot use.
“Why did revenue fall?”	Clarify or provide bounded descriptive analysis if approved.	Avoid unsupported causal claims.

"Export all customer-level data."	Refuse if row-level export is not approved.	Data exposure risk.
"Can I use this in the board pack?"	Explain pilot limitation and escalation route.	Pilot output is not production-approved.

3.5 Support guide template

Support area	What support should know	Escalation route
Access issue	User cannot connect, wrong role, wrong region or missing permission.	Access owner / security lead.
UI issue	User cannot submit question, view answer, provide feedback or load session.	Product / application owner.
Unsupported question	User asks outside approved boundary.	Product owner for backlog classification.
Suspected wrong answer	Answer appears incorrect, misleading or inconsistent.	Evaluation owner + data / semantic owner.
Caveat confusion	User misunderstands limitation, source or answer wording.	Product owner + SME.
Data freshness issue	Data appears stale, incomplete or inconsistent.	Data owner / analytics engineer.
Refusal or clarification issue	System refuses incorrectly or asks poor clarification.	AI / solution architect.
Cost or latency issue	Query or answer is slow, expensive or repeated.	Technical owner / operating owner.
Possible exposure issue	User sees data they should not see or system reveals restricted content.	Security / governance lead immediately.
Incident or pause trigger	Issue may require narrowing or pausing pilot usage.	Pilot decision owner / pause authority.

3.6 Operator and technical documentation checklist

Area	Documentation required	Available?	Location / owner	Gap
MVP version and release notes	Current pilot version, changes since validation, known fixes.			
Architecture overview	High-level flow: UI, model, retrieval, validation, execution, logging.			
Data and metadata scope	Approved assets, metadata sources, caveats and exclusions.			
Access model	User groups, permissions, row / column rules and access lifecycle.			
Logging and traces	Where prompts, retrieval, SQL, validation, answers and feedback are logged.			
Monitoring views	Usage, error, refusal, cost, latency and support dashboards.			
Known limitations	Known technical, data, semantic, model or support limitations.			
Support procedures	Common issue diagnosis and escalation routes.			
Pause and restart route	How affected users, questions, data or the pilot can be paused and restarted.			
Change process	Quick fix, hot fix and material change approval and release process.			
Retest evidence	Required evidence after changes before affected usage resumes.			
Incident record	Where incidents, decisions, owners and follow-up actions are captured.			

3.7 Usage and cost guidance template

Area	Guidance to define
Expected use	What normal pilot use looks like by user or cohort.
Discouraged use	Repeated broad queries, high-volume exploration, unsupported requests, unnecessary refreshes.
Query limits	Any limits on question frequency, concurrency, result size, exports or time windows.
Cost-sensitive behaviour	Question patterns likely to increase model, query or infrastructure cost.
Latency expectations	Typical response time and what users should do when the system is slow.
Heavy-use review	When unusual usage is reviewed and by whom.
Escalation route	Who to contact for slow, expensive, repeated or blocked queries.
Communication	How usage limits or cost-related changes are communicated to users.

Example — illustrative:

Area	Guidance
Expected use	Use the pilot for real sales performance questions during weekly trading review preparation.
Discouraged use	Do not run repeated broad "all regions / all customers / all periods" queries for exploration.
Query limits	Exports and customer-level extracts are not available in the pilot.
Cost-sensitive behaviour	Repeated wide date-range questions may create unnecessary warehouse and model cost.
Heavy-use review	Usage patterns are reviewed weekly by the product owner and operating owner.
Escalation route	Report repeated slow responses through the pilot support channel.

3.8 Onboarding session agenda

- Explain pilot objective and approved boundary.
- Show how to access the pilot and use the UI.
- Walk through supported and unsupported question examples.
- Explain caveats, refusals, clarifications and safe failure.
- Explain responsible use, usage expectations and cost / latency considerations.
- Show how to give feedback and what useful feedback looks like.
- Explain support and escalation route.
- Confirm known limitations and pause conditions.
- Capture user questions and documentation gaps.
- Confirm user is ready to start pilot use.

3.9 Documentation gap log

Gap	Affected audience	Risk	Owner	Decision	Target date
	Users / support / operators			Fix before pilot / accept for pilot / defer / escalate	

Example — illustrative:

Gap	Affected audience	Risk	Owner	Decision	Target date
Caveat wording unclear for current-month revenue.	Users	Users may treat provisional numbers as final.	Product owner	Fix before pilot.	
Support cannot access trace dashboard.	Support	Support depends on build team for diagnosis.	Operating owner	Fix before pilot.	
Cost guidance not included in quick-start note.	Users	Users may create avoidable broad queries.	Product owner	Add before onboarding.	

4 Activity 3: Test pilot access, onboarding, role changes and offboarding

4.1 Pilot access lifecycle test plan

Scenario	Purpose	Test user / role	Expected behaviour	Expected timing	Evidence required	Status
New user onboarding	Confirm approved users can access the pilot.		User can connect and sees only approved scope.		Login, access group, trace sample.	
Existing user access check	Confirm current pilot users still have correct access.		User sees correct data boundary and allowed answer types.		Access test, query / answer trace.	
Role change	Confirm changed role updates access and answer boundary.		System reflects new role and old access no longer applies.		Before / after access evidence.	
Region / business unit change	Confirm row-level boundary changes correctly.		User sees only new approved region / unit.		Query evidence showing applied filter.	
Permission revocation	Confirm restricted capability or data access is removed.		User cannot access removed scope.		Refusal, denied access or blocked query trace.	
Pilot offboarding	Confirm user removed from pilot can no longer access the MVP.		User cannot connect or query.		Access removal evidence.	
Leaver scenario	Confirm joiner / mover / leaver process works where relevant.		Access is removed through normal identity route.		IAM / access log and system check.	

Suggested status values: Planned / Tested / Passed / Passed with caveat / Failed / Not tested / Not applicable.

4.2 Access-change timing record

Scenario	Approval time	Provisioning time	Permission sync time	App / cache refresh time	Total observed time	Expected timing met?	Notes
----------	---------------	-------------------	----------------------	--------------------------	---------------------	----------------------	-------

Add new pilot user						Yes / No / N/A	
Change user role						Yes / No / N/A	
Change region / business unit						Yes / No / N/A	
Revoke permission						Yes / No / N/A	
Remove pilot access						Yes / No / N/A	

Practical note: for a small POC, expected timings may be informal or manually handled. If the pilot is on a path to production, expected and observed timings should be measured and carried into Phase 8 as access-administration evidence.

4.3 Access boundary test matrix

User / group	Approved data scope	Question asked	Expected access behaviour	Actual behaviour	Pass / fail	Evidence
Regional sales manager	Own region only	Revenue by region	Answer only for approved region.			
Sales operations analyst	All regions	Revenue by region	Answer all approved regions.			
Removed pilot user	No access	Any question	Access blocked.			
Changed role user	New approved role only	Prior-scope question	Old scope unavailable.			
Restricted data user	No sensitive data	Customer-level request	Refuse or block query.			

4.4 Traceability evidence checklist

Testing access lifecycle requires enough traceability to confirm which boundary was applied at runtime.

Evidence item	Why it matters	Available?	Location / owner	Restriction
User identifier or pseudonymous user ID	Confirms who asked the question.			
User role / access group at query time	Confirms permission context.			
Question text or redacted question text	Confirms intent and scope.			
Retrieved metadata / policy context	Shows what the system used to ground the answer.			
Generated query or tool call	Confirms table, filter, row limit and access logic.			
Validation result	Shows whether query / access checks passed or failed.			
Final answer or answer summary	Confirms what the user saw.			
Caveats / refusal message	Confirms boundary communication.			
Timestamp and system version	Links result to release, model and policy version.			
Feedback / support ticket link	Links user issue to trace evidence.			

For sensitive data, broad access to raw logs may not be appropriate. Use restricted log access, redacted traces, pseudonymised user identifiers, query summaries or controlled audit views where needed.

4.5 Example access lifecycle test — illustrative

Item	Example
Scenario	Regional sales manager moves from North region to South region.
Expected behaviour	User should no longer see North region results and should only see South region data after permission sync.
Expected timing	Same day after IAM group update and application cache refresh.
Test question before change	"Show revenue for my region last month."
Test question after change	"Show revenue for my region last month."
Evidence required	User role before / after, applied region filter, generated query or query summary, final answer, timestamp.
Pass condition	Answer uses South region filter after change and North region is no longer accessible.
Failure pattern	User continues to see North region data after role change, or answer combines both regions without explanation.

4.6 Access exception and workaround log

Issue	Scenario	Impact	Temporary	Risk	Owner	Decision	Phase 8
-------	----------	--------	-----------	------	-------	----------	---------

			workaround				implication
	Onboarding / role change / offboarding / revocation					Accept / Fix / Constrain / Pause / Escalate	

Example — illustrative:

Issue	Scenario	Impact	Temporary workaround	Risk	Owner	Decision	Phase 8 implication
Role changes require manual app refresh.	Region change	Access update not visible immediately.	Build team refreshes user cache.	Stale permissions during transition.	Technical owner	Accept for pilot with monitoring.	Production needs automated refresh or clear SLA.
Support cannot view redacted query trace.	Access issue	Support cannot diagnose whether issue is access or data.	Escalate to build team.	Slow incident handling.	Operating owner	Fix before broader pilot.	Support trace view required before production.

4.7 Minimum evidence by delivery intent

Delivery intent	Minimum expectation
Small POC / expert test	Manual user provisioning may be acceptable if documented as a shortcut. Basic access checks should still confirm users see only intended data.
Controlled business pilot	Onboarding, access boundary and at least one role / access change scenario should be tested or simulated. Support should know how to escalate access issues.
Production-path MVP	Joiner, mover, leaver, revocation, role-change and permission-sync behaviour should be tested with timing evidence, traceability and named owners.
Sensitive / regulated scope	Access testing should include redacted or restricted trace review, audit evidence, stale-permission checks and clear pause conditions.

4.8 Common mistakes to avoid

Mistake	Why it matters
Testing only initial access	The pilot may work on day one but fail when users change role, leave or need access removed.
Ignoring propagation timing	Access may be “correct eventually” but unsafe during the transition period.
Relying on the build team to provision users	This may be acceptable for a POC, but it is not production evidence unless documented as a shortcut.
No access to traces	Support cannot prove whether the right data boundary, query or refusal was applied.
Logging sensitive data too broadly	Access testing should not create a new exposure risk through unrestricted logs.
Not testing offboarding	Removed users may retain pilot access longer than expected.
No Phase 8 carry-forward	Manual workarounds and timing gaps are forgotten instead of becoming production-readiness items.

5 Activity 4: Run controlled pilot usage and monitored sessions

5.1 Pilot usage plan

Area	Decision / description
Pilot usage objective	What the usage period is intended to test or learn.
User cohort	Approved users or groups participating in the pilot.
User mix	Expert users, normal users, champions, sponsors, occasional users, support users.
Workflow context	Real workflow, meeting, reporting cycle, analyst request, decision process or task.
Usage mode	Self-service use, facilitated session, monitored session or mixed approach.
Expected usage frequency	Daily, weekly, event-driven, reporting-cycle driven or ad hoc.
Pilot duration	Planned window and review points.
Monitoring owner	Person or team reviewing usage and behaviour.
Session facilitator	Person supporting monitored sessions, if used.
Evidence captured	Questions, traces, feedback, observations, support tickets, errors, caveats.
Main caveats	Known limits in user mix, workflow coverage, timing or usage realism.

5.2 User cohort realism check

User type	Included?	Why included / excluded	Evidence caveat
Primary target users			
Expert users			

Normal business users			
Sponsors / champions			
Skeptical or low-adoption users			
Support / operator users			
Users with different access scopes			
Users with different data literacy levels			

Example — illustrative:

User type	Included?	Why included / excluded	Evidence caveat
Sponsors / champions	Yes	They are available and close to the use case.	Positive feedback may overstate broader adoption.
Normal business users	Partially	Two users included from target team.	Limited evidence on wider user behaviour.
Skeptical users	No	Not available during pilot window.	Trust and adoption risks may be under-tested.
Users with different access scopes	Yes	Regional and central users included.	Useful for access-boundary testing.

5.3 Monitored session plan

Session	User / role	Workflow task	Questions to try	Observation focus	Facilitator	Evidence
			Natural questions / prepared task / follow-up questions			Trace / feedback / notes / ticket

Suggested observation focus:

- Does the user know what to ask?
- Does the question fit the approved boundary?
- Does the system clarify, answer, refuse or redirect appropriately?
- Does the user understand the answer and caveats?
- Does the user challenge surprising outputs?
- Does the user know how to give feedback or escalate?
- Does the user rely on the answer appropriately?
- Does the session reveal workflow friction, UI friction or support needs?

5.4 Real-use observation template

Field	Notes
User / role	
Workflow context	
Question asked	
Supported / unsupported / ambiguous	
System response	Answer / clarification / refusal / redirect / error.
User interpretation	Did the user understand the answer, caveat and boundary?
User action	Used answer / challenged answer / abandoned / escalated / asked follow-up.
Friction observed	UI, wording, data, latency, access, caveat, support or workflow issue.
Feedback captured?	Yes / No.
Trace linked?	Yes / No / Not available.
Issue candidate?	Product / data / semantic / access / user education / support / process.
Follow-up action	

5.5 Outside-scope question log

Question	User role	Why outside scope	System behaviour	Correct behaviour?	Backlog implication
		Unsupported metric / data / answer type / access / workflow	Refused / clarified / redirected / answered incorrectly	Yes / No	Ignore / user education / future scope / remediate / pause

Example — illustrative:

Question	User role	Why outside scope	System behaviour	Correct behaviour?	Backlog implication
“What is margin by product?”	Sales manager	Margin not approved for pilot.	Refused and explained boundary.	Yes	Log as future demand.
“Which salesperson performed worst?”	Regional manager	Individual performance excluded.	Clarified then refused.	Yes	Reinforce guidance.
“Forecast revenue next quarter.”	Sponsor	Forecasting not validated.	Gave descriptive trend only.	Partial	Review wording and refusal policy.

5.6 Pilot behaviour evidence log

Behaviour observed	Example / evidence	Risk	Decision
User asks supported questions naturally		Positive workflow-fit signal	Continue / monitor
User repeatedly asks outside scope		Guidance or demand-scope issue	Educate / log / narrow / backlog
User ignores caveats		Over-trust risk	Improve answer format / onboarding
User challenges surprising answer		Good trust calibration	Capture as positive evidence
User cannot connect system output to workflow		Workflow-fit risk	Review use case / user cohort
User needs build-team coaching		Supportability issue	Improve documentation / support route
User abandons due to latency		Adoption and operational risk	Monitor / optimise / constrain
User creates high-cost broad queries		Cost and performance risk	Add usage guidance / limit / monitor

5.7 Facilitator guidance for monitored sessions

Do	Avoid
Ask users to complete realistic workflow tasks.	Steering users only toward questions known to work.
Let users phrase questions naturally.	Rewriting every question for them.
Observe confusion, hesitation and misuse.	Treating confusion as a training inconvenience only.
Capture where users misunderstand caveats or boundaries.	Explaining away every issue during the session.
Link observations to traces, feedback or support tickets.	Relying only on meeting notes.
Record when the user cohort is unusually expert or friendly.	Presenting champion-user evidence as representative adoption.

5.8 Evidence caveat statement template

Use this when the pilot user group is not fully representative.

Pilot findings should be interpreted with caution because the user cohort was weighted towards [sponsors / expert users / motivated champions / a narrow role group]. This provides useful evidence on [workflow value / feasibility / high-intent usage / early adoption], but may under-test [normal user behaviour / low-trust users / wider support load / onboarding friction / misunderstood caveats]. These limitations should be considered before broader rollout or production-readiness decisions.

Example — illustrative: Pilot findings should be interpreted with caution because the user cohort was weighted towards motivated sales operations users and the commercial sponsor. This provides useful evidence on high-intent usage and workflow value, but may under-test normal sales-manager behaviour, onboarding friction and wider support load. These limitations should be considered before broader rollout.

5.9 Session summary template

Item	Summary
Session date	
Users involved	
Workflow tested	
Number of questions asked	
Supported questions	
Unsupported questions	

Clarifications / refusals	
Main positive signals	
Main issues observed	
Trace coverage	Full / partial / missing.
Support issues raised	
Follow-up actions	

5.10 Common mistakes to avoid

Mistake	Why it matters
Using only friendly users	Feedback may overstate adoption, trust and usability.
Running only scripted questions	The pilot does not test real boundary handling or natural user behaviour.
Treating facilitation as normal use	A heavily coached session may hide usability and support issues.
Not logging outside-scope demand	Future scope and user-education needs are lost.
Capturing feedback without traces	The team cannot diagnose whether the issue is user, data, model, access or product.
Ignoring workflow context	A technically correct answer may still not help the real task.
Presenting narrow-user evidence as representative	Phase 8 may underestimate onboarding, support and adoption risk.

6 Activity 5: Capture feedback and monitor pilot signals

6.1 Feedback and monitoring routes

Route	Best use	Watch-outs
In-product feedback	Immediate reaction to a specific answer.	Must capture trace ID or interaction ID.
Support channel	Access issues, incidents, confusion and repeated problems.	Often lacks question / answer context unless requested.
Monitored sessions	Rich qualitative observation of real use.	May overstate usability because users are being observed or coached.
Short user interviews	Workflow fit, trust, unmet needs and adoption barriers.	Should be linked back to actual usage where possible.
Usage-triggered follow-up	Distinctive usage patterns: high, low, unusual or risky use.	Requires usage monitoring and agreed contact route.
Monitoring dashboard	Usage, errors, refusals, cost, latency and support signals.	Only useful if reviewed and linked to decisions.
Review meeting	Pattern review across users, questions, feedback and issues.	Can become anecdotal if not trace-linked.

6.2 Feedback-to-trace and monitoring register

Field	Description
Feedback / signal ID	Unique identifier.
Source	In-product feedback / support / monitored session / dashboard / interview / review meeting.
User / role	User, role or cohort.
Date / time	Timestamp of feedback or signal.
Question type	Supported / unsupported / ambiguous / follow-up / high-cost / access-related.
Interaction / trace ID	Link to trace, log, ticket or session note.
User feedback / observed signal	What the user said or what the system showed.
System behaviour	Answer / clarification / refusal / error / timeout / escalation / high-cost query.
Issue category	Product / data / semantic / access / user education / support / process / model-orchestration / cost-latency.
Severity	Low / medium / high / critical.
Owner	Named owner for review or action.
Decision	Monitor / educate / fix / constrain / backlog / pause / no action.
Next review point	Date, forum or condition.

Suggested issue categories:

- Product or UI issue.
- Data quality, freshness or availability issue.

- Semantic ambiguity or metric-definition issue.
- Access or permission issue.
- Unsupported-scope demand.
- User education or documentation issue.
- Support or operating process issue.
- Model or orchestration issue.
- Cost, latency or performance issue.
- Governance, risk or escalation issue.

6.3 Pilot monitoring dashboard structure

Area	Signal to monitor	Why it matters	Typical owner
Usage	Active users, questions asked, sessions, follow-ups.	Shows adoption and interaction pattern.	Product owner.
User mix	Usage by role, cohort, region or access group.	Shows whether evidence is representative.	Product owner.
Question types	Supported, unsupported, ambiguous, repeated.	Separates value from scope pressure.	Product / evaluation owner.
Trust behaviour	Verification, caveat use, escalation, abandoned sessions.	Shows whether users trust appropriately.	Product / SME.
Errors	Failed queries, timeouts, empty results, UI errors.	Identifies reliability and usability issues.	Technical owner.
Refusals	Refusal volume and refusal appropriateness.	Tests boundary handling.	Evaluation owner.
Clarifications	Clarification frequency and completion.	Tests ambiguity handling.	AI / solution architect.
Support	Tickets, questions, incidents, recurring confusion.	Tests supportability.	Support owner.
Access	Denied access, wrong scope, permission issues.	Tests access lifecycle and exposure risk.	Security / access owner.
Cost	Model cost, query cost, infrastructure cost.	Tests economic sustainability.	Operating owner.
Latency	End-to-end response time and step-level delay.	Tests usability and scale assumptions.	Technical owner.
Risk events	Misleading answer, exposure concern, repeated misuse.	Triggers triage, constraint or pause.	Risk / decision owner.

6.4 Feedback quality checklist

Question	Why it matters	Status
Is the feedback linked to a specific interaction?	Avoids acting on vague comments.	
Is the user role or cohort known?	Helps interpret access, skill level and workflow relevance.	
Is the question type captured?	Separates supported-scope quality from unsupported demand.	
Is the system response available?	Enables diagnosis of answer, caveat, refusal or clarification behaviour.	
Is the generated query or tool call available where needed?	Helps diagnose data, semantic and access issues.	
Is the user's expected answer or outcome known?	Explains whether the issue is correctness, usefulness or expectation mismatch.	
Is the signal linked to monitoring, support or trace evidence?	Prevents feedback and monitoring from becoming separate evidence streams.	
Is there a decision route?	Ensures feedback becomes triage input, not a comment archive.	

6.5 Structured feedback form template

Field	Prompt
User role	What is your role or user group?
Workflow context	What were you trying to do?
Question asked	What did you ask the system?
Answer usefulness	Did the answer help your task? Yes / partly / no.
Answer trust	Did you trust the answer? Why / why not?
Caveat clarity	Were limitations or caveats clear?
Issue type	Wrong / confusing / incomplete / unsupported / too slow / access issue / other.
Expected outcome	What answer or behaviour did you expect?
Action taken	Used answer / verified / escalated / ignored / asked follow-up.
Additional comments	Free text.

6.6 Usage-pattern and distinctive-user follow-up list

Signal	Possible interpretation	Follow-up question
Very high usage	Strong value, curiosity, misuse, broad exploration or support dependency.	What are you using it for, and which outputs affect your workflow?
Very low usage	Poor workflow fit, low trust, access friction, unclear value or weak onboarding.	What stopped you from using it?
Many unsupported questions	Scope mismatch, unclear guidance or demand for future expansion.	Were you aware those questions were out of scope?
Many refusals	Good boundary enforcement or poor scope fit.	Did the refusal help you understand what to do next?
No feedback despite high usage	Feedback friction or overconfidence.	Did anything feel unclear, risky or particularly useful?
Repeated access issues	Lifecycle, permissions or onboarding weakness.	What happened, and how quickly was it resolved?
Long or expensive sessions	Broad queries, unclear user intent or cost-risk pattern.	Were you exploring broadly or trying to answer a specific task?
Frequent follow-ups	Useful exploration or initial answers not clear enough.	Did follow-ups help, or were you compensating for unclear answers?

6.7 Trust and responsible-use indicators

Indicator	Positive signal	Risk signal
Caveat behaviour	Users read and act on caveats.	Users ignore caveats or treat caveated answers as final.
Verification	Users verify surprising outputs.	Users rely on surprising answers without challenge.
Escalation	Users escalate unclear or risky answers.	Users do not know when or how to escalate.
Refusal handling	Users understand why refusal happened.	Users repeatedly try to bypass refusals.
Follow-up use	Users refine questions responsibly.	Users use follow-ups to drift outside scope.
Decision use	Users use answers for approved pilot workflows.	Users use pilot outputs for unapproved reporting or decisions.
Confidence	Users trust useful answers with appropriate caution.	Users either over-trust or avoid the tool entirely.

6.8 Error, refusal and clarification log

Interaction ID	User role	Question type	Event	Expected behaviour?	Severity	Owner	Decision
		Supported / unsupported / ambiguous	Error / refusal / clarification / timeout / empty result	Yes / no / unclear	Low / medium / high / critical		Monitor / fix / educate / constrain / pause

Example — illustrative:

Interaction ID	User role	Question type	Event	Expected behaviour?	Severity	Owner	Decision
T-1042	Regional manager	Margin by product	Refusal	Yes	Low	Product owner	Log as future demand.
T-1087	Sales analyst	Revenue by region	Empty result	No	Medium	Data owner	Investigate data filter.
T-1121	Sponsor	Forecasting	Partial answer	No	High	Evaluation owner	Review refusal rule.

6.9 Cost and latency review template

Metric	Expected range	Observed value	Issue?	Likely cause	Action
Median response time			Yes / no	Model / query / retrieval / UI / network	
95th percentile response time			Yes / no		
Average model calls per interaction			Yes / no		
Average query execution cost			Yes / no		
High-cost interactions			Yes / no	Broad query / repeated follow-up / large result	
Timeout rate			Yes / no	Query complexity / infrastructure / model latency	
Cost per active user			Yes / no	Usage pattern / model choice / query cost	

6.10 High-cost or high-latency interaction review

Interaction	User role	Question	Cost / latency issue	Likely driver	Decision
			High model cost / high query cost / timeout / slow answer	Broad date range / large result / repeated calls / complex join	Optimise / limit / educate / monitor / backlog

Example — illustrative:

Interaction	User role	Question	Cost / latency issue	Likely driver	Decision
T-1204	Sales manager	“Show all customers by product for the last three years.”	Slow and expensive query	Broad result and long time window	Add guidance and result-size limit.
T-1218	Sales analyst	“Compare every region, segment and month.”	High query cost	Multi-dimensional broad query	Review query guardrail.

6.11 Positive-value evidence template

Positive feedback should be linked to real workflow value, not only satisfaction.

User / role	Workflow	Question / interaction	Value observed	Evidence	Caveat
			Faster answer / better exploration / reduced analyst dependency / improved consistency / better decision support	Trace / feedback / interview / session note	

Example — illustrative:

User / role	Workflow	Question / interaction	Value observed	Evidence	Caveat
Regional sales manager	Weekly trading review	Revenue and pipeline conversion by segment.	Reduced preparation time and enabled follow-up questions during review prep.	Trace IDs, user feedback and monitored session note.	User is a motivated pilot champion; broader adoption not yet proven.

6.12 Feedback and monitoring diagnosis guide

User says / signal shows...	Do not assume...	Check...
“The answer is wrong.”	The model failed.	Data source, metric definition, filter, access boundary, generated query and user expectation.
“I do not trust it.”	The answer is wrong.	Caveat clarity, source visibility, prior user experience and sensitivity of workflow.
“It refused my question.”	Refusal is a bug.	Whether the question was unsupported, ambiguous, unsafe or outside access rights.
“It is too slow.”	The UI is slow.	Model calls, query execution, retrieval, validation, network and result size.
“This is useful.”	Production value is proven.	Which workflow it helped, frequency, risk level and whether the user is representative.
Low usage	The use case has no value.	Onboarding, trust, timing, workflow fit, access friction and competing tools.
High usage	The system is ready to scale.	Value, misuse, cost, latency, support burden and unsupported demand.
Repeated unsupported questions	The product is failing.	Boundary handling, user guidance and future-scope demand.

6.13 Feedback and monitoring review cadence

Cadence	Review focus	Participants
Daily or near-daily during early pilot	Blocking issues, access problems, misleading answers, exposure risks, major latency or support load.	Product, support, technical owner, security where needed.
Weekly	Usage patterns, trust, refusals, support load, cost, unsupported demand, quick fixes and backlog.	Product, business SME, support, evaluation, data / semantic owner.
After hot fix	Did the fix work and avoid regression?	Product, technical owner, support, affected risk owner.
End-of-pilot	Value, trust, workflow fit, supportability, risk, Phase 8 backlog and decision recommendation.	Decision forum and accountable owners.

6.14 Pilot signal summary template

Area	Summary
Feedback volume	Number of feedback items by source.
Trace linkage	Share of material feedback linked to trace / interaction evidence.
Usage	Active users, sessions, question volume, follow-ups.

Representative use	Whether usage came from normal users, experts, sponsors or champions.
Value pull	Evidence of sustained use, requests for more scope or workflow integration.
Trust behaviour	Caveat use, verification, over-trust, low trust, escalation.
Boundary handling	Unsupported questions, refusals, clarifications, scope pressure.
Quality / errors	Errors, empty results, failed queries, misleading answers.
Supportability	Support volume, repeated issues, escalation needs, build-team dependency.
Cost and latency	Cost drivers, slow interactions, heavy users, broad queries.
Risk events	Exposure concerns, misuse, material answer issues, pause triggers.
Backlog implications	Fix now, constrain, Phase 8, later enhancement, no action.
Decision implication	Continue, fix, constrain, pause, extend or Phase 8 backlog.

6.15 Common mistakes to avoid

Mistake	Why it matters
Relying on satisfaction scores	Users may like a tool they misuse or over-trust.
Measuring usage volume only	High usage may hide misuse, cost pressure or unsupported demand.
Ignoring low usage	Low usage may reveal poor workflow fit, weak onboarding or low trust.
Collecting feedback without trace IDs	The team cannot diagnose what actually happened.
Treating unsupported demand as product failure	Scope expansion becomes uncontrolled.
Treating representative demand as production readiness	Value pull matters, but safety and supportability still need Phase 8.
Counting refusals without reviewing them	A refusal may be correct, confusing or blocking valid use.
Monitoring cost too late	Cost patterns can become normalised before ownership is agreed.
Not separating users by cohort	Champion-user usage may not represent wider adoption.
Dashboards without review cadence	Monitoring only matters if it drives decisions.
Letting feedback bypass triage	Issues become anecdotes instead of controlled pilot evidence.

7 Activity 6: Review answer quality, caveat understanding and user interpretation

7.1 Real-interaction answer-quality review template

Interaction ID	User role	Question type	Supported?	Answer quality	Caveat handling	User interpretation	Decision
		Supported / ambiguous / unsupported	Yes / No / Partial	Correct / partial / wrong / unclear	Clear / weak / missing / ignored	Understood / misunderstood / unknown	Accept / fix / educate / interview / constrain / pause

Suggested review dimensions:

- Correct data source.
- Correct metric definition.
- Correct filters, grain and date logic.
- Correct access boundary.
- Appropriate caveats.
- Appropriate refusal or clarification.
- No unsupported causal claim, recommendation or confidence.
- User understood the answer in the intended way.

7.2 Answer-quality scoring guide

Score	Meaning	Typical action
Pass	Answer is correct, grounded, caveated and understood.	Keep as positive evidence.
Pass with caveat	Answer is mostly correct but wording, caveat or interpretation risk exists.	Improve wording or guidance.

Partial fail	Answer is technically usable but incomplete, unclear or weakly caveated.	Fix, retest or monitor.
Fail	Answer is wrong, misleading, unsafe or misunderstood in a material way.	Triage, constrain or pause affected scope.
Unknown	Trace is insufficient or user interpretation is unclear.	Follow up through interview or support review.

7.3 Caveat and limitation review

Caveat area	Review question	Finding	Action
Visibility	Was the caveat visible enough to the user?		
Wording	Was the caveat understandable to a non-expert user?		
Timing	Was the caveat shown before the user relied on the answer?		
Specificity	Did it explain the actual limitation, not a generic warning?		
User action	Did the user change behaviour because of the caveat?		
Escalation	Did the caveat explain when to verify or escalate?		

Example — illustrative:

Caveat area	Review question	Finding	Action
Visibility	Was the caveat visible enough?	Caveat appeared after the answer and was not noticed.	Move caveat closer to the key number.
Specificity	Did it explain the limitation?	“Data may be incomplete” was too generic.	State that current-month revenue is provisional.
User action	Did the user act on it?	User used provisional value as final.	Improve wording and onboarding example.

7.4 Misunderstood-answer log

Interaction ID	Answer was technically correct?	What user appeared to understand	Intended meaning	Risk	Follow-up needed?	Decision
	Yes / No / Partial			Low / Medium / High	Yes / No	Wording fix / user education / interview / constraint / no action

Example — illustrative:

Interaction ID	Answer was technically correct?	What user appeared to understand	Intended meaning	Risk	Follow-up needed?	Decision
T-2041	Yes	User treated month-to-date revenue as final month revenue.	Current month is provisional and incomplete.	High	Yes	Change answer label and caveat placement.
T-2098	Yes	User read a variance as causal explanation.	System only showed descriptive change.	Medium	Yes	Add wording: “This does not explain cause.”

7.5 Follow-up interview trigger list

Some interactions cannot be judged from traces alone. Use this list to identify cases requiring a short interview.

Trigger	Why interview is needed	Example question to ask user
User gave negative feedback but trace looks correct	Need to understand expectation mismatch.	What answer did you expect, and why?
User accepted a caveated answer	Need to know whether caveat changed behaviour.	How did the caveat affect your decision?
User asked many follow-ups	Need to know whether exploration was useful or compensating for unclear answers.	Were the follow-ups helping, or was the first answer unclear?
User abandoned session	Need to understand whether issue was trust, UX, latency or workflow fit.	What made you stop using the answer?
User used answer in workflow	Need to confirm decision context and risk.	How did you use the answer afterwards?
User repeatedly asked unsupported questions	Need to separate guidance gap from unmet demand.	Did you know this was outside scope?
User did not give feedback despite heavy use	Need to surface silent issues or overconfidence.	Did anything feel unclear, risky or particularly useful?

7.6 User interpretation interview guide

Area	Questions
------	-----------

Context	What were you trying to achieve when you asked the question?
Expectation	What answer did you expect before using the system?
Interpretation	What did you understand from the answer?
Caveats	Did you notice the caveat or limitation? What did it mean to you?
Trust	Did you trust the answer? Why or why not?
Verification	Did you check the answer elsewhere or ask someone?
Action	Did the answer change what you did next?
Confusion	Was anything unclear, missing or misleading?
Boundary	Did you understand whether the question was in scope?
Improvement	What would make the answer safer or more useful?

7.7 Refusal and clarification understanding review

Interaction ID	Behaviour	Expected?	User understood why?	Issue	Action
	Refusal / clarification / redirect	Yes / No / Partial	Yes / No / Unknown	Boundary unclear / wording unclear / refusal wrong / user expectation mismatch	Educate / rewrite / fix / backlog

Example — illustrative:

Interaction ID	Behaviour	Expected?	User understood why?	Issue	Action
T-2120	Refusal	Yes	No	User thought system was broken, not out of scope.	Improve refusal wording and quick-start guide.
T-2152	Clarification	Yes	Yes	User selected region and continued.	Positive evidence.

7.8 Answer-format improvement log

Issue	Example	Risk	Proposed improvement	Owner	Retest needed?
Caveat too low in answer		User misses limitation.	Move caveat near key number.		Yes / No
Metric label unclear		User misunderstands definition.	Add metric label and definition link.		Yes / No
Descriptive answer sounds causal		User infers unsupported cause.	Add “descriptive only” wording.		Yes / No
Refusal too generic		User does not know what to ask instead.	Add boundary and alternative suggestion.		Yes / No
Confidence wording misleading		User over-trusts answer.	Use plain-language confidence / evidence statement.		Yes / No

7.9 Real-interaction review summary

Area	Summary
Interactions reviewed	Number and selection method.
Supported questions	Pass / partial / fail patterns.
Unsupported questions	Refusal, redirect and boundary handling.
Ambiguous questions	Clarification quality and completion.
Caveat understanding	Main caveats understood or misunderstood.
Misunderstood correct answers	Count, examples and risk level.
Follow-up interviews needed	Users / interactions selected and reason.
Main fixes	Wording, UI, caveat, guidance, data, semantic or model changes.
Pilot decision implication	Continue / fix / constrain / pause / Phase 8 backlog.

7.10 Common mistakes to avoid

Mistake	Why it matters
Reviewing answer correctness only	A correct answer may still be misunderstood or misused.
Assuming caveats are understood because they are displayed	Users may not notice, read or act on them.

Treating refusals as failures by default	A refusal may be the correct safe behaviour.
Ignoring technically correct but misleading wording	The answer can create unsupported causal or decision claims.
Diagnosing user intent from logs alone	Some interactions require short interviews.
Reviewing only bad answers	Positive interactions show what wording and patterns should be reused.
Not linking interpretation issues to traces	The team cannot determine whether the problem is user, product, data or model behaviour.

8 Activity 7: Triage issues, manage backlog and controlled change

8.1 Pilot issue triage and backlog log

Issue ID	Source	Description	User / cohort	Question type	Evidence link	Issue category	Severity	Decision	Owner	Target phase	Status
	Feedback / trace / monitoring / support / access test / answer review			Supported / unsupported / ambiguous	Trace / ticket / session note	Product / data / semantic / access / model / user education / support / process / scope / cost	Low / medium / high / critical	Continue / monitor / educate / fix / constrain / pause / redesign / stop		Pilot / Phase 8 / later	

Suggested issue categories:

- Product or UI issue.
- Data quality, freshness or availability issue.
- Semantic issue: metric, definition, grain, join, filter or caveat.
- Access or permission issue.
- Model or orchestration issue.
- Unsupported-scope demand.
- User education or documentation issue.
- Support or operating process issue.
- Cost, latency or performance issue.
- Governance, risk or escalation issue.

8.2 Severity and urgency guide

Severity should reflect both technical impact and business context.

Level	Typical meaning	Example response
Low	Minor issue with limited impact and no material risk.	Monitor, document or fix later.
Medium	Repeated issue, user confusion or friction affecting pilot quality.	Fix, educate, update guidance or backlog.
High	Issue affects supported answers, trust, access, supportability or decision use.	Triage urgently; consider constraint or pause.
Critical	Data exposure, materially misleading answer, unsafe reliance, broken auditability or major control failure.	Pause affected scope immediately and escalate.

Urgency modifiers:

Modifier	Why it changes priority
Senior decision maker affected	Small issues may damage trust or block sponsorship.
Critical workflow timing	Issue appears during board, trading, reporting or operational cycle.
Repeated occurrence	Pattern suggests systemic failure, not isolated bug.
High-confidence user impact	User relied on, or was likely to rely on, the output.
Exposure or access risk	Even one case may justify immediate pause.
Public / visible pilot moment	Trust impact may exceed technical severity.
Easy safe fix available	Quick fix may be justified if tested and controlled.

8.3 Decision route matrix

Decision	Use when	Evidence required	Typical owner
Continue	Issue is low risk and understood.	Issue note and monitoring signal.	Product owner.
Monitor	Pattern is unclear or not yet material.	Trace / feedback and review point.	Product / evaluation owner.
Educate	Issue is mainly user guidance or misunderstanding.	User feedback, examples, guidance gap.	Product owner / business SME.
Fix during pilot	Issue is bounded, low-risk and testable.	Fix note, test evidence, communication plan.	Product / technical owner.
Constrain	Risk is localised to users, questions, data or answer types.	Affected scope and constraint rationale.	Product owner + risk owner.
Pause affected scope	Continued use may be unsafe or misleading.	Severity evidence, pause trigger, restart condition.	Pause authority.
Redesign	Issue reveals structural workflow, data, semantic, orchestration or support weakness.	Root-cause evidence and impact assessment.	Decision forum.
Stop	Value, risk or operating model is not credible enough to continue.	Pilot evidence pack and sponsor decision.	Sponsor / decision forum.

8.4 Root-cause classification guide

Symptom	Possible root causes to test
User says answer is wrong	Data source, metric definition, filter, access boundary, query generation, user expectation.
User asks unsupported questions	Guidance gap, natural demand beyond scope, wrong cohort, future backlog opportunity.
Caveats ignored	Poor caveat placement, unclear wording, over-trust, weak onboarding, high-pressure workflow.
Refusal frustrates user	Wrong refusal, unclear boundary, genuine unmet need, poor redirect, unsupported scope.
Slow interaction	Model latency, query cost, retrieval, validation, result size, infrastructure or network.
Support cannot diagnose issue	Missing traces, unclear ownership, weak runbook, no support access, unclear issue taxonomy.
Access issue	IAM sync, role mapping, row-level security, stale permissions, app cache, manual provisioning.
Positive feedback but risky use	Over-trust, wrong decision context, sponsor bias, unclear pilot limitation.
Repeated hot fixes	Structural product, data, semantic, access or workflow issue being patched rather than solved.

8.5 Continue / constrain / pause / restart decision record

Field	Decision
Issue ID	
Decision	Continue / constrain / pause / restart / redesign / stop
Affected users	
Affected question types	
Affected data / assets	
Evidence reviewed	Trace, feedback, ticket, monitoring, answer review, access test.
Rationale	
Immediate action	
Communication required	Users / support / sponsor / governance / security.
Restart condition	
Owner	
Review date	

Example — illustrative:

Field	Decision
Issue ID	P7-034
Decision	Constrain
Affected users	Regional sales managers
Affected question types	Month-to-date revenue comparison
Evidence reviewed	Three traces, two feedback comments, caveat review.
Rationale	Answers were technically correct but users interpreted provisional current-month values as final.

Immediate action	Add stronger caveat and restrict use in reporting-cycle pack.
Communication required	Notify pilot users and support channel.
Restart condition	Retest answer wording and confirm caveat understanding with two users.
Owner	Product owner
Review date	

8.6 Hot-fix and change classification guide

Item type	Definition	Typical route
Hot fix	Urgent bounded fix needed to keep pilot safe or usable.	Fix, test, communicate, record.
Quick win	Low-risk improvement that improves usability, guidance or supportability.	Fix if inside approved boundary.
Pilot constraint	Temporary reduction of users, questions, data or answer types.	Approve and communicate.
User guidance update	Clarify supported questions, caveats, feedback route or cost guidance.	Product review and user communication.
Support guidance update	Add triage route, known issue or escalation step.	Support owner review.
UI wording fix	Improve labels, caveat placement or refusal text.	Product review and light regression test.
Answer-format fix	Change answer structure, source display or caveat layout.	Answer-quality regression test.
Prompt / model-routing change	Adjust prompt, routing or model task allocation.	Evaluation / regression test before release.
Validation-rule change	Change query validation, refusal, guardrail or limit.	Technical and risk review; retest failure cases.
Metadata / semantic change	Change metric description, synonym, caveat or join guidance.	Semantic owner approval and regression test.
Access handling change	Role mapping, permission logic, row-level rule or user group update.	Security / access review and access tests.
Scope change	New users, data, questions, answer types or domains.	Formal pilot boundary approval.
Phase 8 item	Required before production readiness or broader release.	Add to production-readiness backlog.
Later enhancement	Useful but not required for pilot or production readiness.	Add to future roadmap.
Stop / redesign signal	Issue challenges value, architecture, workflow fit or operating model.	Escalate to decision forum.

8.7 Hot-fix decision checklist

Question	Yes / no	Notes
Is the issue inside the approved pilot boundary?		
Is the fix bounded and well understood?		
Is the risk of the change lower than the risk of leaving the issue unresolved?		
Can the fix be tested before release?		
Are relevant regression tests available?		
Does the fix affect answer behaviour, caveats, refusal logic, access, model routing or supported scope?		
Does the fix require user or support communication?		
Is there a rollback or reversal route if the fix creates regression?		
Has the owner approved release into the pilot?		

Recommended decision values: Release / Release with caveat / Defer / Constrain / Pause / Escalate.

8.8 Controlled change / release record

Field	Description
Fix / change ID	
Issue addressed	
Change type	Guidance / UI / prompt / validation / metadata / access / monitoring / support / other.
Reason for change	Why this cannot wait for Phase 8 or later backlog.
Affected users / scope	
Expected behaviour before change	
Expected behaviour after change	
Tests completed	
Regression checks completed	
Communication required	Users / support / operators / sponsor / governance.

Release owner	
Approval	
Release date	
Rollback / reversal route	
Post-release review date	

Example — illustrative:

Field	Description
Fix / change ID	HF-007
Issue addressed	Users misunderstood current-month revenue caveat.
Change type	Answer format / caveat wording.
Reason for change	Supported answer was technically correct but creating interpretation risk.
Affected users / scope	Sales performance pilot users.
Expected behaviour after change	Caveat appears next to current-month number and states “provisional month-to-date value”.
Tests completed	Golden question rerun, caveat review, answer-format check.
Communication required	Pilot users and support channel.
Release owner	Product owner.
Post-release review date	Next pilot review meeting.

8.9 Regression check template

Regression area	Check	Required for
Supported questions	Previously supported questions still answer correctly.	Prompt, model, metadata, answer-format changes.
Unsupported questions	Refusals and redirects still work.	Prompt, validation, refusal wording changes.
Access boundary	User permissions and data restrictions still apply.	Access, metadata, query or execution changes.
Caveat handling	Required caveats still appear and are understandable.	Answer-format, prompt or metadata changes.
Query safety	Queries remain bounded, valid and within approved assets.	Prompt, SQL generation, validation changes.
Cost / latency	Change does not materially increase response time or cost.	Model, retrieval, query or validation changes.
Traceability	Logs still capture question, context, query, answer and feedback.	UI, orchestration, logging or release changes.
User experience	Change does not confuse users or contradict guidance.	UI, wording, caveat or answer-format changes.

8.10 User and support change note template

Use this when a change affects user-visible behaviour.

Section	Content
What changed	Plain-language description of the change.
Why it changed	Issue, feedback or risk addressed.
Who is affected	Users, roles, workflows or question types affected.
What users should do differently	Any change in behaviour, interpretation or escalation.
Known caveats	Remaining limitation after the fix.
Support note	What support should know if users ask about the change.
Effective date	When the change is live.
Contact / escalation	Where to report issues.

Example — illustrative:

We have updated how current-month revenue answers are displayed. Current-month values are now labelled as provisional month-to-date figures and the caveat appears next to the number. This change was made because some users interpreted current-month values as final monthly results. Please continue to verify provisional figures before using them in reporting-cycle decisions.

8.11 Pilot change log

Change ID	Date	Change summary	Change type	Owner	Tests completed	Users informed?	Status
			Hot fix / quick win / constraint / material			Yes / no / N/A	Draft / released / rolled back /

		change				deferred
--	--	--------	--	--	--	----------

8.12 Phase 8 carry-forward and backlog separation

Backlog item	Source evidence	Why not fully resolved in Phase 7	Production-readiness implication	Owner	Priority	Required before go-live?
	Trace / feedback / support / monitoring / interview / test	Accepted for pilot / workaround / outside pilot scope / needs design			High / medium / low	Yes / no / conditional

Example — illustrative:

Backlog item	Why not fully resolved in Phase 7	Production-readiness implication	Owner	Priority
Role-change propagation timing not automated	Manual workaround accepted for pilot.	Phase 8 must define access-change SLA and automation route.	Operating owner	High
Support trace view is limited	Build team can support small pilot.	Broader release requires redacted trace dashboard for support.	Technical owner	High
Margin demand frequent but unsupported	Metric not approved.	Future scope requires semantic definition and validation.	Semantic owner	Medium

8.13 Triage and change review agenda

- Review new high-severity issues and pause triggers.
- Review issues linked to supported questions first.
- Review access, exposure, auditability and misleading-answer risks.
- Review repeated user confusion, unsupported demand and support burden.
- Classify root cause and severity.
- Decide continue, monitor, educate, fix, constrain, pause, redesign or stop.
- Confirm whether the item is hot fix, quick win, pilot constraint, Phase 8 item, later enhancement or redesign signal.
- Confirm regression checks and communication required for any pilot change.
- Assign owner, evidence required and review date.
- Update pilot backlog and Phase 8 implications.

8.14 Example triage and change cases — illustrative

Case	Classification	Decision	Rationale
Senior sponsor sees incorrect regional filter in supported answer.	Access / query / trust issue	Pause affected question type	Small number of users affected, but high trust and exposure impact.
Users repeatedly ask for margin, which is out of scope.	Unsupported demand / semantic gap	Continue, log demand, reinforce guidance	Not a defect if refusal is correct; useful future-scope signal.
Caveat displayed but users miss it.	Product / user interpretation	Quick fix	Move caveat closer to key number and retest.
High latency for broad queries.	Cost / performance	Constrain and educate	Add result-size guidance and monitor expensive patterns.
Support cannot access trace evidence.	Support / observability	Fix before expansion	Pilot can continue for small cohort, but Phase 8 needs support trace access.
Low usage from target users.	Adoption / workflow fit	Extend and interview users	Insufficient evidence; may be timing, onboarding or value issue.
Hot fix changes refusal wording.	Controlled change	Release with regression test and user note	Behaviour change affects user understanding and support guidance.

8.15 Common mistakes to avoid

Mistake	Why it matters
Treating all feedback equally	A minor issue in a critical workflow may be more urgent than a larger issue in an unused path.
Fixing before classifying	The team may patch symptoms while leaving data, semantic or access causes unresolved.

Fixing severe emergency	The team may patch symptoms while leaving data, semantic or access issues unresolved.
Ignoring user role and decision context	Severity depends on who saw the issue and what they were trying to do.
Using backlog as a dumping ground	Some issues require immediate constraint, pause or escalation.
Treating unsupported demand as defect	Scope pressure should be separated from approved-scope performance.
Continuing despite broken boundary	The pilot stops being controlled if unsafe exposure continues.
Not recording judgement calls	Triage is partly judgement-based; decisions need rationale, owner and review point.
Releasing hot fixes without telling users	Users may misinterpret changed behaviour or compare old and new answers incorrectly.
Treating a hot fix as a small change without regression tests	Small wording, prompt or validation changes can alter answer behaviour materially.
Expanding scope through fixes	The pilot boundary changes without approval.
Fixing symptoms repeatedly	The team may miss a structural data, semantic, access or workflow issue.
No version history	The team cannot explain which behaviour users saw at which point in the pilot.
No rollback route	A failed fix can make the pilot less safe than before.

9 Activity 8: Assess adoption, workflow fit, supportability and Phase 8 implications

9.1 Pilot decision summary template

Area	Summary
Pilot objective	
Pilot users / cohort	
Pilot duration	
Evidence coverage	Usage, feedback, traces, monitored sessions, support tickets, interviews.
Main value signals	
Main risk signals	
Supportability findings	
Production-readiness gaps	
Recommended decision	Proceed to Phase 8 / proceed with constraints / extend / narrow / pause / redesign / stop.
Decision rationale	
Decision owner	
Review date	

9.2 Adoption evidence review

Signal	Evidence	Interpretation	Caveat
Sustained use		Strong workflow value / curiosity / dependency / other	
Repeat users		Users are returning after initial trial	
Low usage		Poor fit / weak onboarding / low trust / wrong timing	
Requests for more scope		Value pull / scope pressure / unmet need	
Requests for integration		Workflow relevance and potential operating value	
Champion-heavy usage		Useful early signal but not representative	
Representative-user usage		Stronger adoption evidence	
Drop-off after first use		Curiosity spike or disappointment	

Example — illustrative:

Signal	Evidence	Interpretation	Caveat
Sustained use by regional managers	8 of 12 target users used the tool weekly for trading review prep.	Strong workflow-fit signal.	Pilot users were motivated and supported closely.
Requests for margin questions	15 unsupported requests across 6 users.	Clear demand for scope expansion.	Margin metric not approved; not evidence that production scope should expand now.

9.3 Workflow-fit assessment

Workflow	User group	Current pain point	Pilot evidence	Fit assessment	Decision implication
----------	------------	--------------------	----------------	----------------	----------------------

		Slow answer / analyst dependency / dashboard gap / inconsistent metric / exploration need		Strong / partial / weak / unclear	Proceed / improve / narrow / stop
--	--	---	--	-----------------------------------	-----------------------------------

Suggested fit indicators:

- Users ask real workflow questions, not only test prompts.
- The MVP supports follow-up exploration better than existing dashboards.
- Users can use answers without build-team coaching.
- The answer is timely enough for the decision cycle.
- The tool reduces friction without creating unacceptable support or trust risk.
- Users ask for workflow integration, not just broader novelty access.

9.4 Value evidence template

Value area	Evidence to capture	Example
Speed	Time saved versus current process.	Faster trading review preparation.
Exploration	Follow-up questions answered without analyst queue.	User drills from region to segment.
Consistency	Users receive answers using approved definitions.	Revenue definition applied consistently.
Trust	Users understand source, caveat and limitation.	User verifies surprising result before using.
Analyst leverage	Reduced dependency on repetitive analyst requests.	Analyst focuses on exceptions, not simple cuts.
Decision support	Better preparation, faster review or clearer issue identification.	Manager identifies declining segment earlier.

9.5 Supportability assessment

Area	Evidence	Assessment	Phase 8 implication
Support volume	Tickets, questions, incidents	Low / manageable / high / unclear	
Support type	Access, UI, answer, caveat, latency, data, unsupported demand		
Escalation need	Share requiring data, AI, security or build team		
Build-team dependency	How often build team had to explain, fix or diagnose		
Access lifecycle	Onboarding, role change, offboarding, revocation		
Monitoring usability	Could support / operators use dashboards and traces?		
Documentation quality	Could users and support self-serve common questions?		
Incident handling	Were incidents triaged, owned and resolved?		

9.6 Phase 8 production-readiness backlog

Backlog item	Source evidence	Why it matters for production	Owner	Priority	Required before go-live?
	Trace / feedback / support / monitoring / interview / test			High / Medium / Low	Yes / No / Conditional

Suggested backlog categories:

- Access lifecycle and administration.
- Monitoring and alert thresholds.
- Support model and runbook.
- Documentation and onboarding.
- Cost ownership and budget controls.
- Evaluation and regression tests.
- Security / privacy / audit evidence.
- Reliability, latency and resilience.

- Data / semantic remediation.
- UX, answer format and caveat clarity.
- Scope expansion candidates.
- Governance and release decisioning.

9.7 Decision recommendation guide

Recommendation	Use when
Proceed to Phase 8	Value, controlled use, supportability and risk evidence are strong enough.
Proceed with constraints	Phase 8 can start, but with reduced scope, users, data, answer types or assumptions.
Extend pilot	Evidence is promising but incomplete, unrepresentative or affected by timing.
Narrow pilot	Value exists, but risk or support burden is concentrated in part of the scope.
Pause and remediate	Continued use may mislead users, expose data, overload support or damage trust.
Redesign	Workflow, data, semantic, orchestration, access or support assumptions are structurally weak.
Stop	The value case, risk profile or operating model is not credible enough to continue.

9.8 Decision wording examples — illustrative

Proceed to Phase 8

The pilot provides sufficient evidence to proceed to production-readiness assessment for the approved sales performance scope. Representative users used the MVP in weekly workflow preparation, feedback was trace-linked, support load was manageable, and material risks are captured in the Phase 8 backlog. Production readiness should focus on access lifecycle, monitoring thresholds, support model, documentation and run-cost ownership.

Proceed with constraints

The pilot should proceed to Phase 8 with constraints. Revenue and pipeline questions showed clear workflow value, but current-month caveats and unsupported margin demand create interpretation and scope risk. Phase 8 should exclude margin, strengthen caveat presentation, define production support ownership and validate access lifecycle before release.

Extend pilot

The pilot should be extended because evidence is promising but insufficient. Usage was concentrated among sponsors and expert users, with limited evidence from representative business users. The extension should focus on normal-user onboarding, real workflow use, support burden and trust calibration.

Pause and remediate

The affected pilot scope should pause while remediation is completed. Supported answers were technically correct but repeatedly misunderstood in a reporting-cycle context. Usage should restart only after answer wording, caveat placement, user guidance and regression evidence are reviewed.

Stop

The pilot should stop for this use case. Usage remained low after onboarding and follow-up, users did not connect the MVP to a recurring workflow, and the support and semantic effort required is not justified by the observed value. A dashboard or analyst workflow improvement should be reconsidered.

9.9 Pilot caveat statement template

Use this where the evidence is useful but not fully representative.

Pilot findings should be interpreted with caution because [user cohort / timing / workflow coverage / data scope / support model] was limited. The pilot provides useful evidence on [value / feasibility / user behaviour / supportability], but under-tests [normal users / low-adoption users / broader access lifecycle / high-volume usage / production support / regulated workflow]. These caveats should be reflected in the Phase 8 scope and decision.

Example — illustrative:

Pilot findings should be interpreted with caution because usage was concentrated among motivated sales operations users and the commercial sponsor. The pilot provides useful evidence on workflow value and high-intent usage, but under-tests normal sales-manager onboarding, low-trust users and broader support load. These caveats should be reflected in the Phase

8 scope and decision.

9.10 Common mistakes to avoid

Mistake	Why it matters
Treating positive feedback as production evidence	Users may like the tool before it is safe, supportable or cost-controlled.
Ignoring representative-user evidence	Sustained use by real target users is one of the strongest value signals.
Presenting champion usage as broad adoption	Sponsor enthusiasm may overstate normal-user readiness.
Mixing blockers and enhancements	Phase 8 needs to know what blocks production, not just what users want.
Ignoring support dependency	A pilot that relies on the build team may not scale operationally.
Expanding scope because users ask for more	Demand is useful evidence, but new scope still needs semantic, access and validation work.
Forgetting evidence caveats	Phase 8 may overestimate readiness if pilot cohort, timing or workflow coverage was narrow.
Ending with a feature backlog only	The pilot should produce a production-readiness decision, not just product requests.